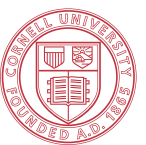

CS5112: Algorithms and Data Structures for Applications

Data representations and real-world efficiency

Ramin Zabih

Sources: Wikipedia; Kevin Wayne, [Kleinberg/Tardos](#)



Administrivia

- Prelim (midterm) date: Wednesday October 16
 - In class, closed book
 - Review session in class on October 7
- HW1 is out, due by 10/7
- Lectures will be recorded “Real Soon Now”
- Minor RDZ error in lecture Wednesday

Lecture Outline

- Memory and data representations
- Memory allocation
- The memory hierarchy
- Practical implications for data structure/algorithm design

Computers and arrays

- Notion of “constant time” operations relies on a hardware model
- Computer memory is a large array
 - We will call it M , elements $M[0]$, $M[1]$, etc. Integers, for simplicity.
- In constant time, a computer can:
 - Read any element of M
 - Change any element of M to another element
 - Perform any simple arithmetic operation
- This is **more or less** what the hardware manual for an x86 describes


Data structures in memory

- Today we will focus on implementing various data structures using the array M
- We have a contract (API) we wish to fulfill
 - Example: stack contract
 - If we push X onto S and then pop S , we get back X , and S is as before
- We want to fulfill this contract using the memory array M
- We'll look at related data structures also

Linked lists


- You can add items to it, and you can process the items you added
- Unlike a stack, don't need to pop it in order to get at all the items
- Adding items takes constant time
- Getting each item takes linear time
 - Example: searching the list for an element
- You can also delete an element
 - This is **always** the hard part of any data structure, and source of most bugs
 - If you don't believe me, look up delete on red-black trees

Linked lists



Linked lists as memory arrays

- We'll implement linked lists using M
 - This is very close to what the hardware does





- A linked list contains “cells”
- A value, and where the next cell is
 - We will represent cells by a pair of adjacent array entries

A few details

- I will draw odd numbered entries in blue and even ones in red
 - Odd entries are values
 - Number interpreted as list elements
 - Even ones are “cells”
 - Number interpreted as index of the next cell
 - AKA *location*, *address*, or ***pointer***
- The first cell is M(1) and M(2), for now
- The last cell has 0, i.e. pointer to M(0)
 - Also called a “null pointer”

Example



Access every item in a list

- Start at the first cell, $M(1)$ and $M(2)$
- Access the first value, $M(1)$
- The next cell is at location $v=M(2)$
- If $v = 0$, we are done
- Otherwise, access the next value, $M(v)$
- The next cell is at location $v' = M(v+1)$
- Keep on going until the next cell is at location k , such that $M(k+1) = 0$

Adding a header

- We can represent the linked list just by the initial cell, but this is problematic
 - Think about deleting the last cell!
 - More subtly, it's also a problem for insertion
- Instead, we will have a few entries in M that are not cells, but instead hold some information about the list
 - Specifically, a pointer to the first element, we use $M[1]$
 - Having a counter is also useful at times, we use $M[2]$

Insert and delete

- To insert an element, create a new cell and splice it into the list
 - Various places it could go
 - Need to update header info, and the cell (if any) before the new one
- To delete an element we simply need to update the header and to change some pointers
 - Basically, need to “skip over” deleted element

Linked list with header

- Initial list:

1	2	3	4	5	6	7	8	9
5	2	2	0	1	3	X	X	X

- Insert (end):

1	2	3	4	5	6	7	8	9
5	3	2	7	1	3	E	0	0

- Insert (start):

1	2	3	4	5	6	7	8	9
7	3	2	0	1	3	S	5	0

- Delete (end):

1	2	3	4	5	6	7	8	9
5	1	2	0	1	0	X	X	X



- Delete (start):

1	2	3	4	5	6	7	8	9
3	1	2	0	1	3	X	X	X

Linked lists and dequeues


- Linked lists support insertion/deletion at beginning or end
 - What is the complexity of these 4 operations?
 - Is this a good way to implement queues?
- How do we modify linked lists so that we can delete from the end in constant time?
 - We clearly need a pointer to the last element in the header
 - But this is not enough!

Doubly linked lists



A doubly-linked list in memory

- M[1] points to start
- M[2] points to end
- M[3] is number of cells
- Cells have 3 elements
 - Before, value, after
 - Before & after are pointers
 - Can be null



1	2	3	4	5	6	7	8	9
4	7	2	0	8	7	4	4	0

Notes on doubly-linked lists

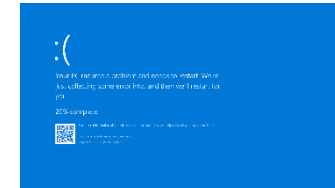
- Inserting and deleting is fast, but very easy to get wrong
 - Try it on all cases, especially trivial ones
 - Look for **invariants**: statements that must be true of any valid list
 - Debug your code by checking invariants
 - In C/C++, this is done via *assert*
 - Most languages have a facility like this built in
 - But if not, you can just write your own!

Memory allocation

- We've assumed the hardware supplied us with a huge array M
 - When we need more storage, we just grab locations at the end
 - Keep track of next free memory location
 - What can go wrong?
 - Consider repeatedly adding, deleting an item
- Notice that when we delete items from a linked list we simply changed pointers so that the items were inaccessible
 - But, they still waste space!

Storage reclamation

- Need to figure out certain locations can be re-used (“garbage”)
 - If too conservative, programs will run slower and slower
 - “memory leak”
 - If it’s too aggressive, your program will crash
 - [BSOD](#) or [Sad Mac](#)
- Suppose your linked list was corrupted
 - Digression: Why do computers crash when we read/write an arbitrary location? Must they??
 - See: 3131961357
 - 0xbaadf00d



Two basic options

- Computer keeps track of free storage
- Manual storage reclamation
 - Programmer has to explicitly free storage
 - Languages: C, C++, assembler
- Automatic storage reclamation
 - Computer will free storage for you
 - Languages: Matlab, Java, C#, Scheme, SML, Python
- Basic tradeoff is program speed versus programmer effort
 - The computer isn't (usually!) as smart as a programmer

Allocation issues

- Computer keeps a linked list of free storage blocks (“freelist”)
 - For each block, keep size and location
 - When asked for new storage, get from freelist
- Surprisingly important question:
 - Which block do you supply?
- Different sized blocks
 - Actually, several different lists (for efficiency)
 - “Buddy block” system, see: <http://www.memorymanagement.org/>

Manual storage reclamation

- Programmers always ask for a block of memory of a certain size
 - In C, explicitly declare when it is free
- Desirable but complex invariants:
 - Everything should be freed when it is no longer going to be used
 - And, it should be freed exactly once!
 - Also want to minimize fragmentation
- Serious C programmer write their own memory allocator!

Automatic storage reclamation

- First challenge: figure out what memory locations are in use by the programmer
 - They are the ones that the programmer can get to (i.e., in a linked list currently in use)
 - Anything that has a name the programmer can get to
 - Start with something that has a name, then chase pointers
 - Some programs can access everything...
 - These languages don't allow such programs!

How to garbage collect?

- Need to distinguish a value from a pointer
 - Garbage collector doesn't know anything about the programmer's data representation
- Usual solution: use highest-order bit
 - The pointer to location 3 is represented as $2^n + 3$, where n is usually 31 (word size – 1)
 - Some computers had hardware support to tell pointers from values
 - Nice idea that died out in the 90's

Simple algorithm: mark-sweep

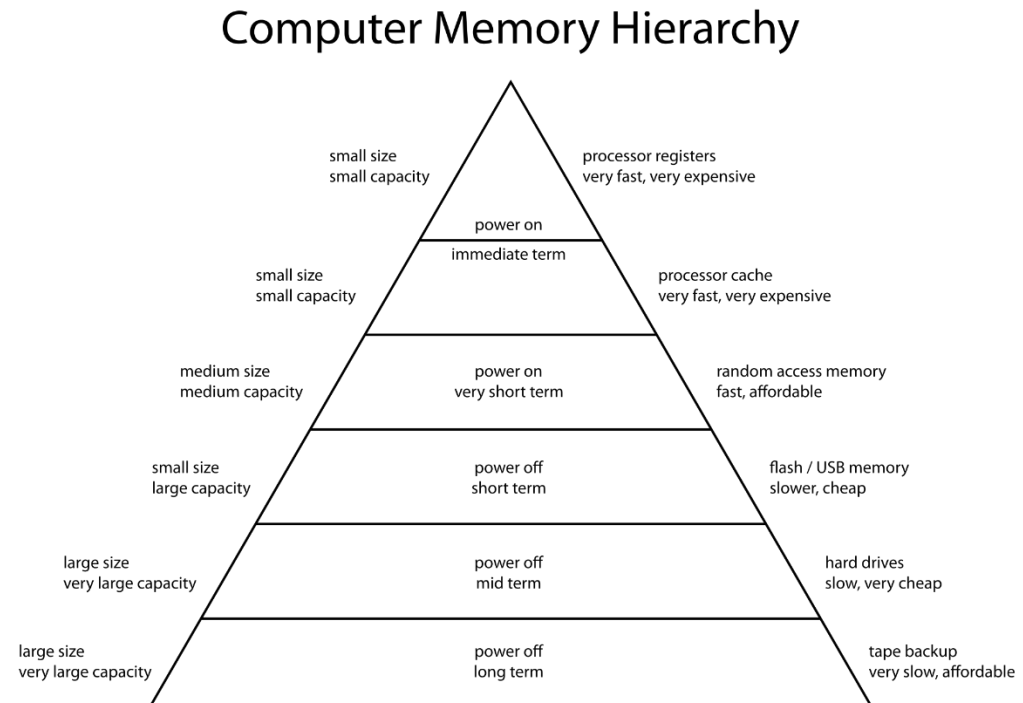
- Chase the pointers from the root set until you are done
 - Modifying everything as you go (“mark”)
 - Typically, set a bit
- Everything not marked is garbage, goes back on the free list
- Problem: typically a lot more garbage than non-garbage, and need to examine all the garbage with mark-sweep
- Python uses reference counting, or more complex schemes

Speed of memory access

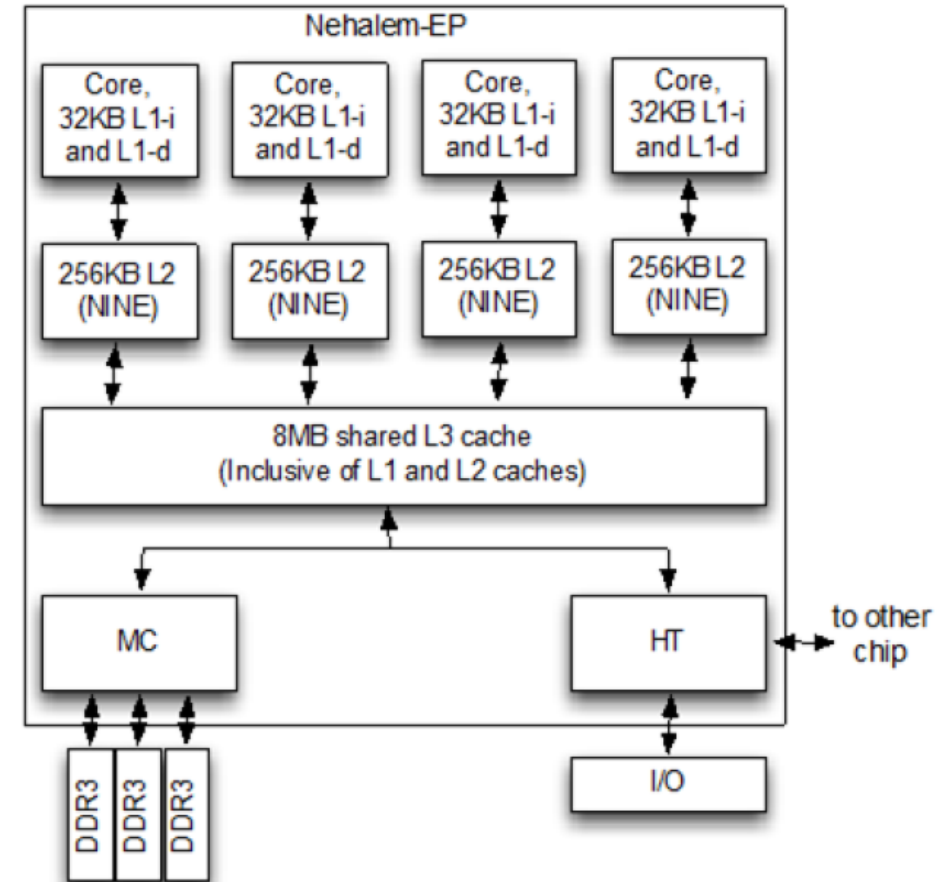
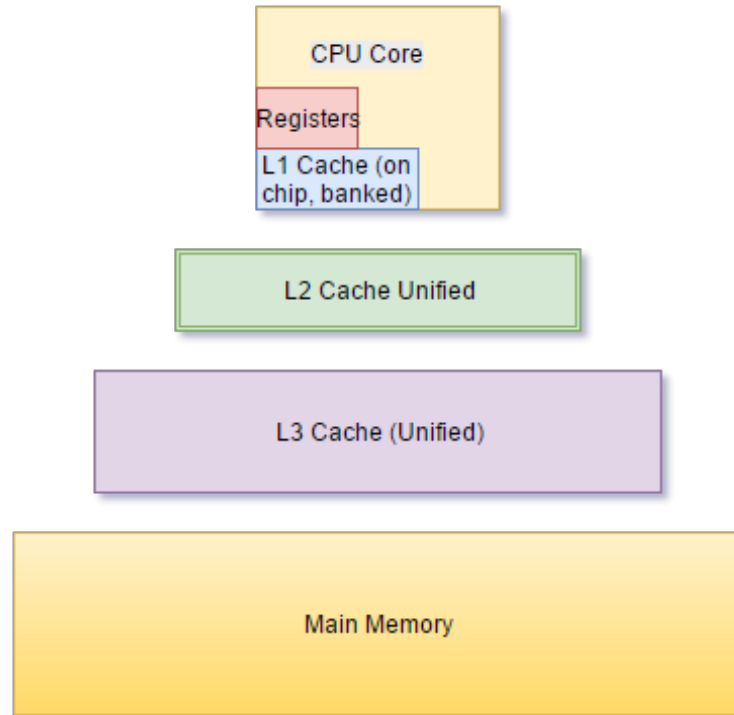
- We made the simplifying assumption that it's constant
 - It's not, and this matters. A lot.
 - Physics dictates memory is small and fast vs large and slow
- There are huge differences in speed depending on how a program accesses memory
 - Certain serious programmers obsess over this (e.g. video games)

Memory hierarchy

- Basic idea: look for data in small fast memory first
- On a miss, look in slower bigger memory
 - Write is similar to read



Real example: Intel Nehalem



- Example timings (not real):
 - main memory = 50 ns, L1 = 1 ns (10% miss rate), L2 = 5 ns (1% miss rate), L3 = 10 ns (0.2% miss rate) [[Wikipedia](#)]

Practical consequences

- Data structures and algorithms that are smart about memory access can be WAY faster than ones that are dumb
- Key idea: exploit locality in space and time
- Consequence of cache layout, and refill strategy
 - Classical example: array access
 - Row major or column major order?
 - Prim versus Kruskal?
- Cache management is quite complex on modern hardware