

CS5112: Algorithms and Data Structures for Applications

Association rules

Ramin Zabih

Some content from: Wikipedia/Google image search; Harrington;
J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmids.org>



Administrivia

- Prelim (midterm) date scheduled for Wednesday October 16
 - This would ruin your holiday weekend, so we're delaying it a week
 - In class, closed book on Wednesday October 23
 - Review session in class on October 21
- HW1 is due by 10/7
 - We will be more forceful about groups in the future
- Lectures will be recorded “Real Soon Now”

Lecture Outline

- Unsupervised learning and association rule mining
- Some logical identities
- Useful ways to think about (discrete) probabilities
- Frequent item set data mining
- The Apriori algorithm

Learning from big data

- The term is ill-defined, but lots of important examples
 - Fraud detection in credit card transactions
 - Loyalty programs
 - Predicting consumer preferences
 - And exploiting them?
- Primarily rely on labeled data
 - “Supervised learning”
- Labeled data is **expensive**

Unsupervised learning

- What can we learn in the absence of a labeled data set?
- Main unsupervised areas are:
 - Clustering (e.g., k-means)
 - Low dimensional structure (not covered in CS5112)
 - Associations (today's lecture)
- Digression: semi-supervised learning
 - Use clusters to propagate a sparse set of labels

Useful logical identities

- Consider true/false propositions p, q, r, \dots
- The below can be proved by, e.g. truth tables

$$(p \Rightarrow q) \equiv (\neg p \vee q) \equiv (\neg q \Rightarrow \neg p)$$

$$(p \wedge q \Rightarrow r) \equiv (p \wedge \neg r \Rightarrow \neg q)$$

$$(p \Rightarrow q \wedge r) \vdash (p \Rightarrow q)$$

Example transactions

<i>TID</i>	<i>Items</i>
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

- Rule discovered: Coke→Diaper

Things can go badly wrong...





Association rules

- Learn rules that are supported by your data
- Rules are co-occurrence, not causality!
 - Very clear in the propositional formulation (symmetry)
- Beer and diapers legend
 - What do you do with an association rule?
- In practice you don't want too many of them
 - Need to act on them

Support and confidence

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

- Have both a computational and probabilistic interpretation
- **Support** of an itemset is the percentage of the transactions containing that itemset
 - In our example, support of Milk is 
 - Support of a rule is the support of LHS
 - Not all papers use this definition, sometimes it's the support of LHS \cup RHS
- **Confidence** of an association rule is percentage of transactions where that rule is correct
 - Confidence of Milk \rightarrow Bread is 

Probabilistic view

- “The basket contains beer” can be viewed as a proposition p , or as a 0/1 random variable
- Consider rule: p generally follows from $q \wedge r$
- Can be viewed as the idea that $P(p|q, r)$ is large
 - Support is joint probability $P(p, q, r)$
- Confidence is conditional probability $P(p|q, r)$
 - Note that $P(p, q, r) = P(p|q, r)P(q, r)$
 - Symmetry: $P(p, q, r) = P(r|p, q)P(p, q) = P(q|p, r)P(p, r)$

Possible worlds

	Coffee	¬Coffee	
Tea	15	5	20
¬Tea	75	5	80
	90	10	100

- Each possible world (= outcome) has some probability
 - This assigns probabilities to subset of possible worlds (= events)
 - Just add up their probabilities
 - Note: works well in the discrete case (only!)
- Joint probability $P(p, q, r)$ is the probability of the possible worlds where all 3 propositions are true
- Conditional probability $P(p|q, r)$ is the probability of the possible worlds where p is true
 - Normalized by only considering possible worlds where q, r are true

Association rule learning

- All rules with support $\geq s$ and confidence $\geq c$
- We focus on finding sets with large support (why?)
 - Called **frequent (item) sets**
- Many rules from same item set, different c

<i>TID</i>	<i>Items</i>
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

$\{\text{Milk, Diaper}\} \rightarrow \{\text{Beer}\} (s=0.6, c=0.67)$

$\{\text{Milk, Beer}\} \rightarrow \{\text{Diaper}\} (s=0.4, c=1.0)$

$\{\text{Diaper, Beer}\} \rightarrow \{\text{Milk}\} (s=0.6, c=0.67)$

$\{\text{Beer}\} \rightarrow \{\text{Milk, Diaper}\} (s=0.6, c=0.67)$

$\{\text{Diaper}\} \rightarrow \{\text{Milk, Beer}\} (s=0.8, c=0.5)$

$\{\text{Milk}\} \rightarrow \{\text{Diaper, Beer}\} (s=0.8, c=0.5)$

Beyond confidence

- Sometimes other measures are useful
- Motivating example:

	Coffee	\neg Coffee	
Tea	15	5	20
\neg Tea	75	5	80
	90	10	100


- $c = P(\text{Coffee}|\text{Tea}) = 0.75$
 - But $P(\text{Coffee}) = 0.9$
 - And $P(\text{Coffee}|\neg\text{Tea}) = 0.9375$
- **Lift** is one solution: $\frac{P(\text{Coffee}|\text{Tea})}{P(\text{Coffee})} = \frac{0.75}{0.9} < 1$

Lift and identities

- Formal definition: $\text{lift}(p \rightarrow q) = \frac{P(q|p)}{P(q)}$
 - Intuitively: how much more or less common is q in the possible worlds where p is true, compared to in general
 - Does buying p make you more or less likely to buy q ?
 - When $\text{lift}=1$ the two items being purchased are independent
 - When $\text{lift}<1$ the items are substitutes
 - When $\text{lift}>1$ then the items complement each other
- Note: $\text{lift}(p \rightarrow q) = \text{lift}(q \rightarrow p)$ (because: algebra)

PB&J example

- Item set is $\{P, J, B\}$
- Consider the rule $\{P, J\} \rightarrow B$
- Support of 0.03 for LHS means P, J in 3% of transactions
- Confidence of 0.82 for rule means 82% of transactions that purchase P, J also purchase B
- If B had support of 43% then the rule has a lift of



Fields of sets

- Consider a set with n elements
- We can arrange all of its 2^n subsets into a lattice
 - Via union and intersection
- This structure is called a field of sets

Example




Figure 11.2 All possible itemsets from the available set {0, 1, 2, 3}

The A Priori Principle

- Problem: exponentially many item sets
- As we grow an item set, its support goes down
- If an item set is frequent, all of its subsets are frequent
 - How could beer and diapers be popular, if beer were not popular?
- If an item set is infrequent, all of its supersets are infrequent
 - If beer and coke is infrequent, so is beer and coke and diapers

Example




Figure 11.3 All possible itemsets shown, with infrequent itemsets shaded in gray. With the knowledge that the set {2,3} is infrequent, we can deduce that {0,2,3}, {1,2,3}, and {0,1,2,3} are also infrequent, and we don't need to compute their support.

Apriori algorithm

- Given a support threshold and a set of transactions
- Find frequent single items
- To go from frequent k -tuples to (possibly) frequent $k + 1$ tuples, combine k -tuples with frequent single items
 - Ex: from 2-tuples to 3-tuples
 - Combining a frequent single item and a frequent 2-tuple can lead to an infrequent 3-tuple
- Stop when no more frequent tuples

From frequent item sets to rules

- In bricks and mortar situations, usually require about 1% support and 50% confidence
- Given a frequent item set with k elements, there are $k - 1$ logically equivalent rules
 - Of the form $p_1 \wedge p_2 \wedge \cdots p_{k-1} \Rightarrow p_k$
- We know that the LHS is frequent, so we can easily calculate the confidence of this rule

Apriori plus and minus

- Plus: Fast, runs on huge data sets, easy to interpret
- Rules with high confidence but low support are missed
 - Classic example: vodka → caviar
- Very widely cited paper, many algorithmic variants
 - Ex: a transaction with no frequent itemsets can be ignored later
 - Including, for example, dynamic programming!