


Distributed Computing

Map/Reduce

- Developed at Google, many other implementations since (Apache Hadoop, etc.)
- Designed for processing big data sets in parallel
- Framework provides
 - Parallelization
 - Redundancy
 - Fault-tolerance
 - Speed...?*
- Abstracts the split-apply-combine analysis strategy
- Works for a multi-threaded environment or a networked cluster environment

Map/Reduce

- Master/Controller node reads input and feeds (k_1, v_1) pairs to worker nodes (mappers)
- Map
 - Apply a user-provided map function on data keyed by k_1 and generates a list of outputs
 - $(k_1, v_1) \rightarrow [(k_2, v_2)]$
- Shuffle
 - Redistributes map outputs so that similarly key'd data are on the same shard
 - $hash(k_2) \% N_R$
 - $[(k_2, v_2)] \rightarrow (k_2, [v_2])$
- Reduce
 - Apply a user-provided reduce function to combine the list of v_2 values into a final output
 - $(k_2, [v_2]) \rightarrow [v_3]$
- Master/Controller node collates reducer outputs into a single final output



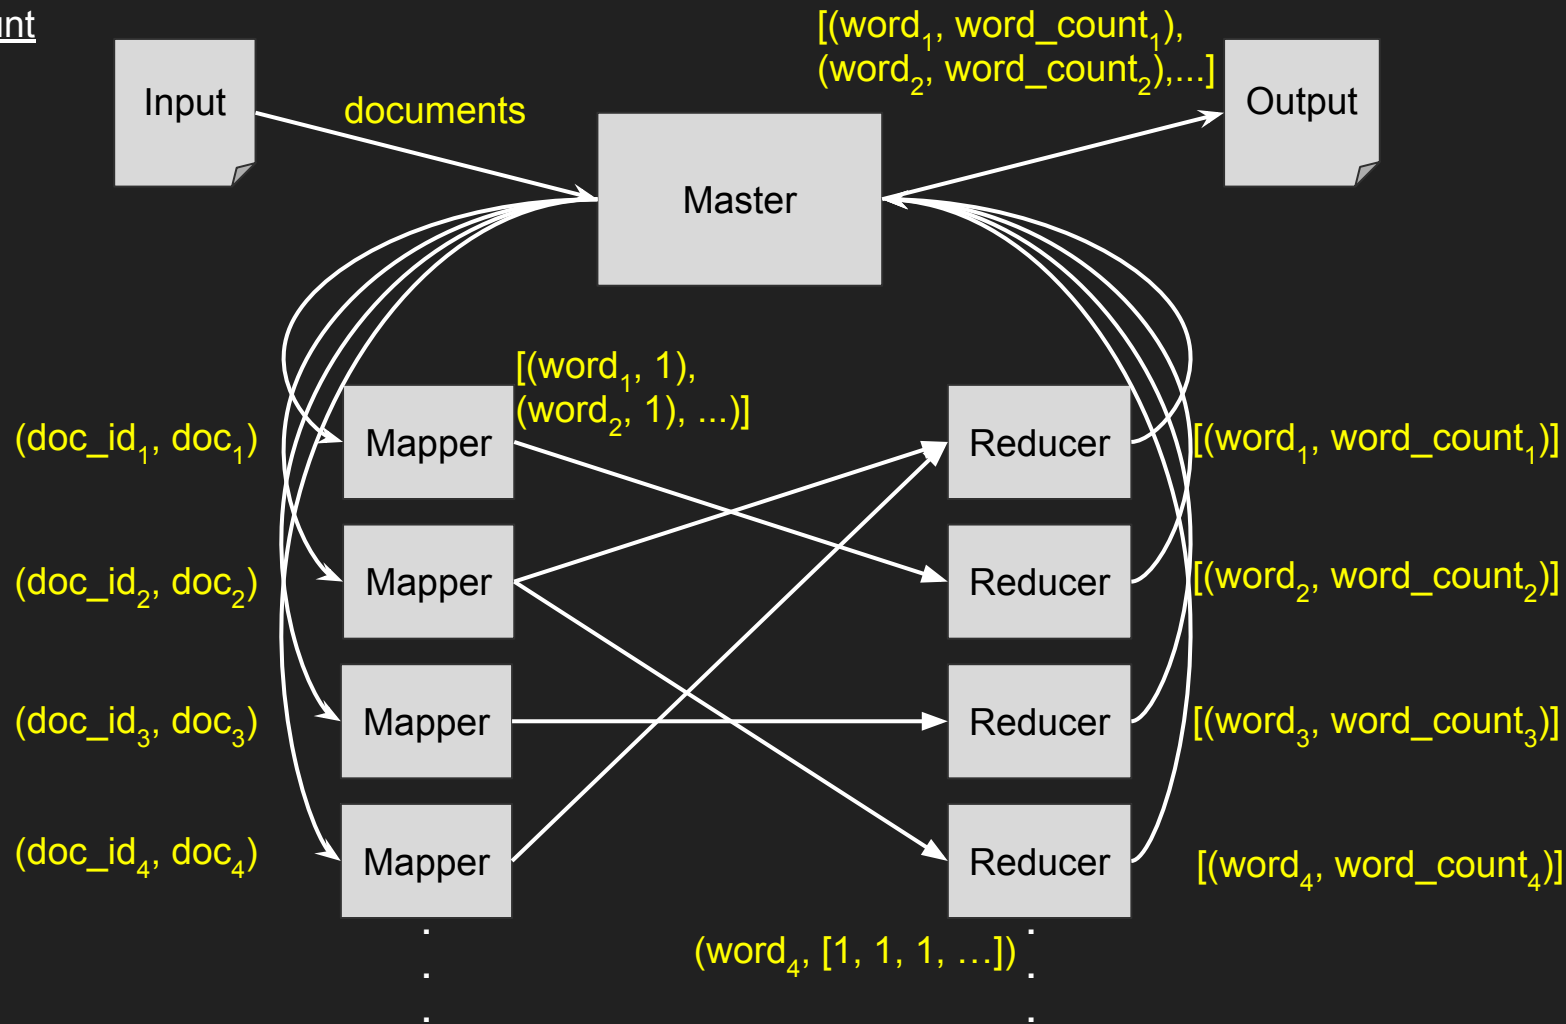
Map/Reduce

- Framework is essentially the same for every application
 - User only needs to supply the Map and Reduce functions
- What happens if a Mapper dies?
 - Controller keeps track of keys that have been successfully mapped
 - Can always send key to a new mapper
- What happens if Reducer dies?
 - Controller keeps track of the keys that have been successfully reduced
 - Mappers save their results on disk in the shard, so any other machine can start from there
- Not necessarily fast
 - “Shuffle” stage actually hides some complication about combining everything
 - Network/communication latency also a factor

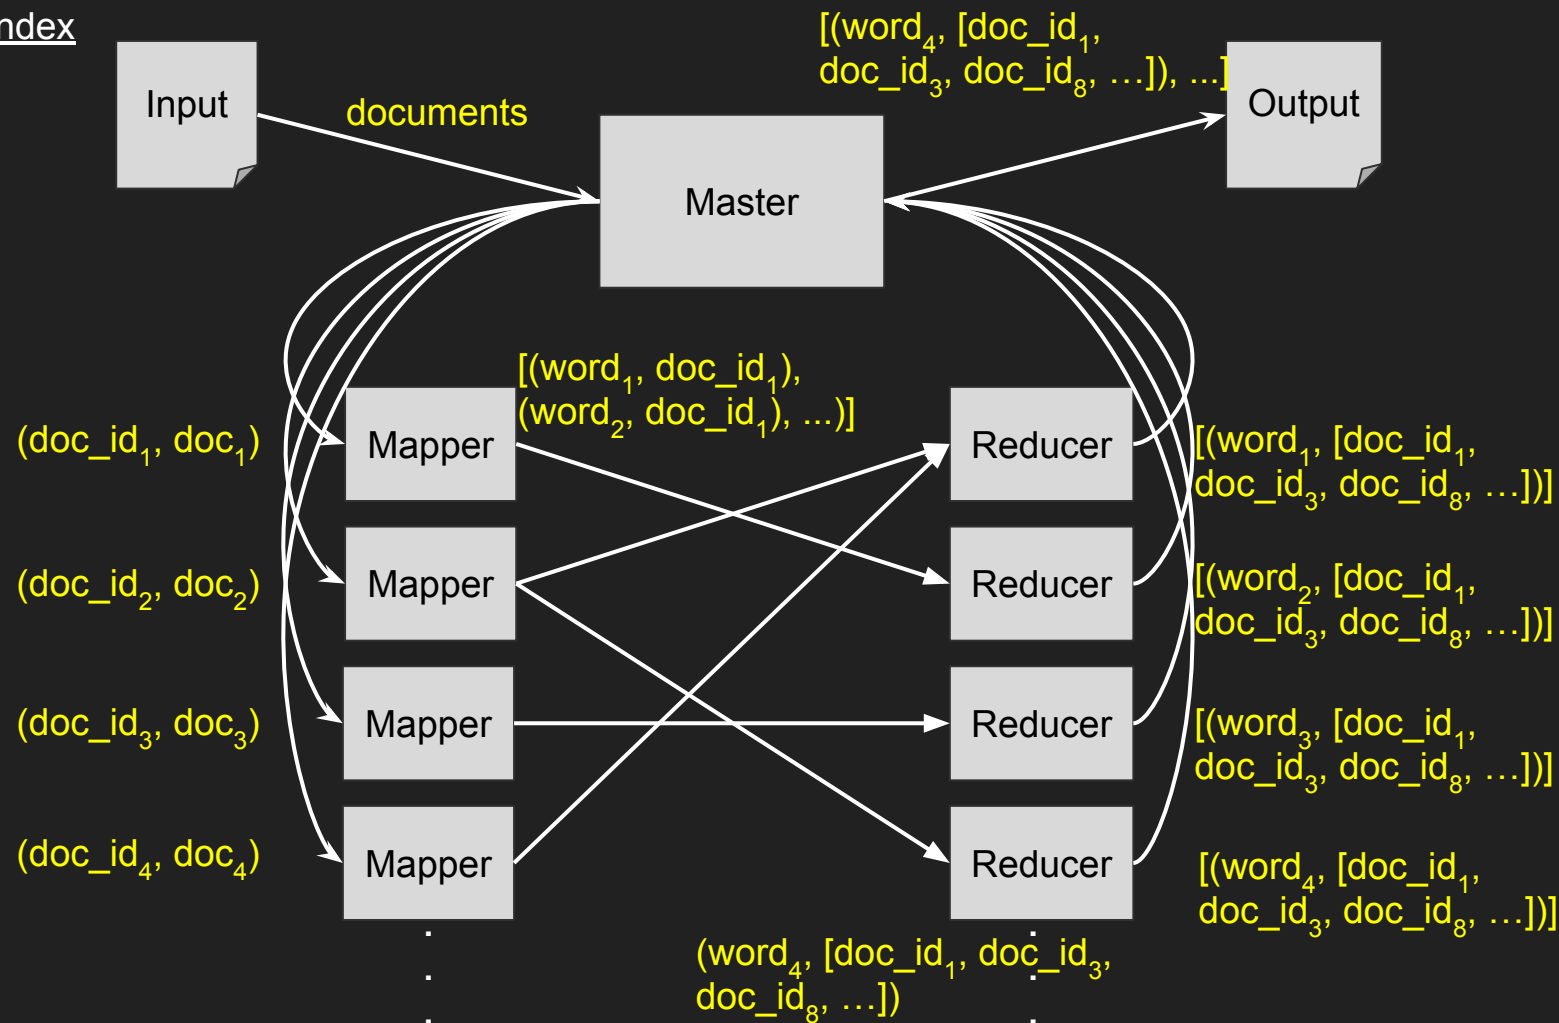
Map/Reduce Example Applications

- Word-count
 - Given a list of documents, for every word in any document output how many times each word appears across all documents
- Inverted Index
 - Given a list of documents, for every word in any document output for each word which documents those words appear in
- *k*-means clustering
 - Given n points in space, find k clusters that best-fit the data set
 - Specifically word_count vectors for documents; using word counts to cluster similar documents
 - Spoilers: requires multiple map/reduce runs!
- PageRank
 - Given a graph of web links, find the probability that a user will end up on any given page
 - Spoilers: requires multiple map/reduce runs!

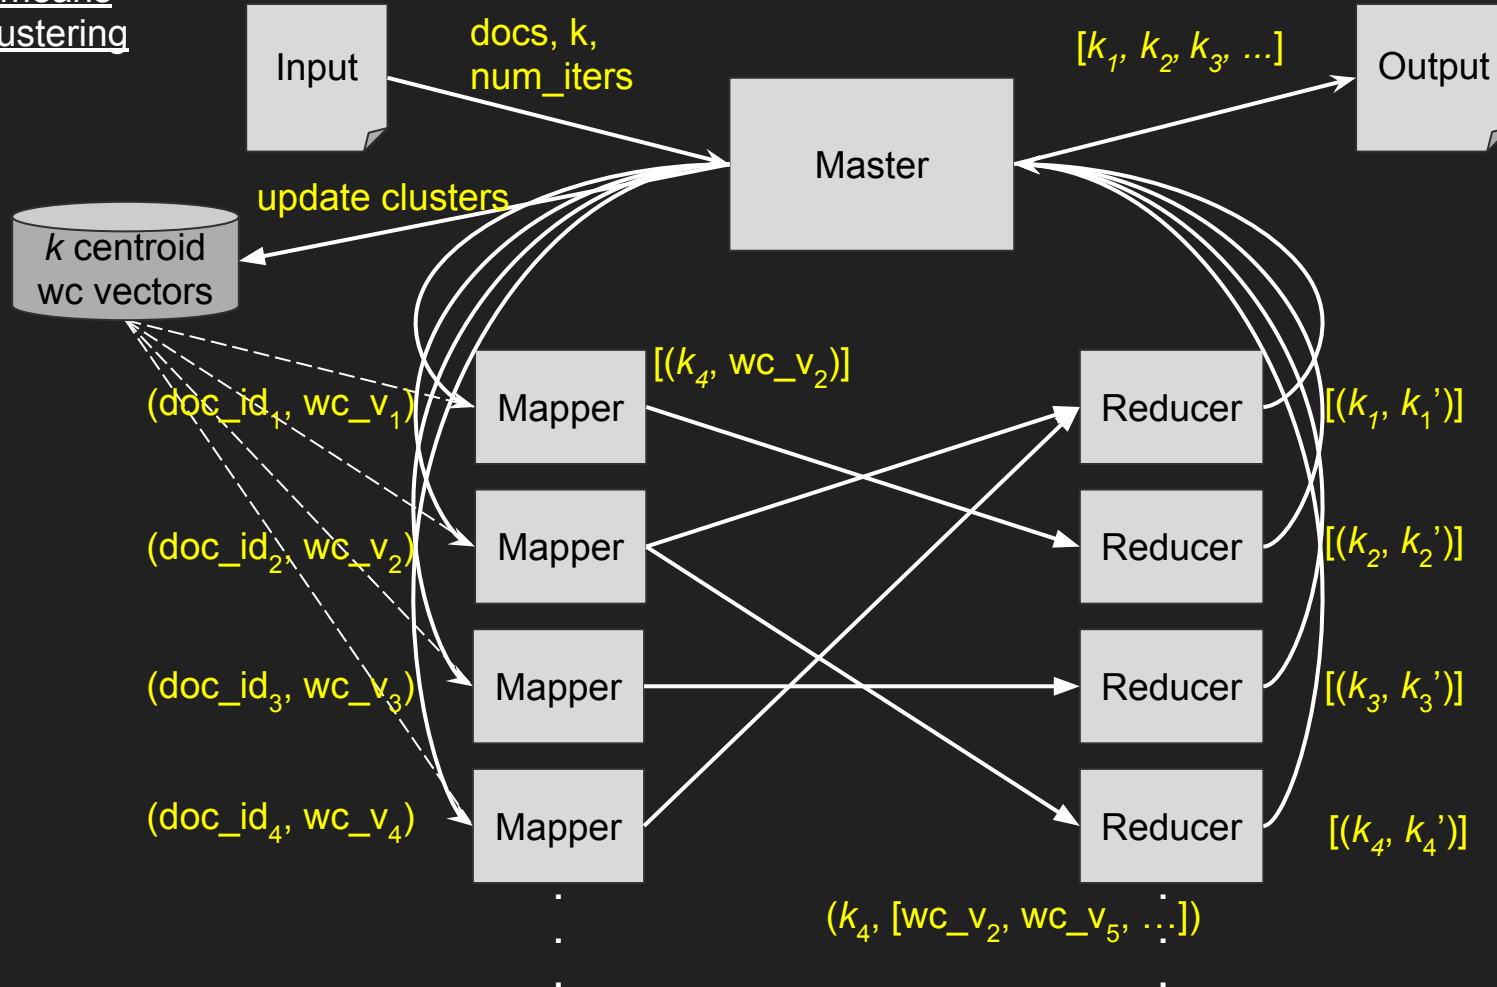
Word Count



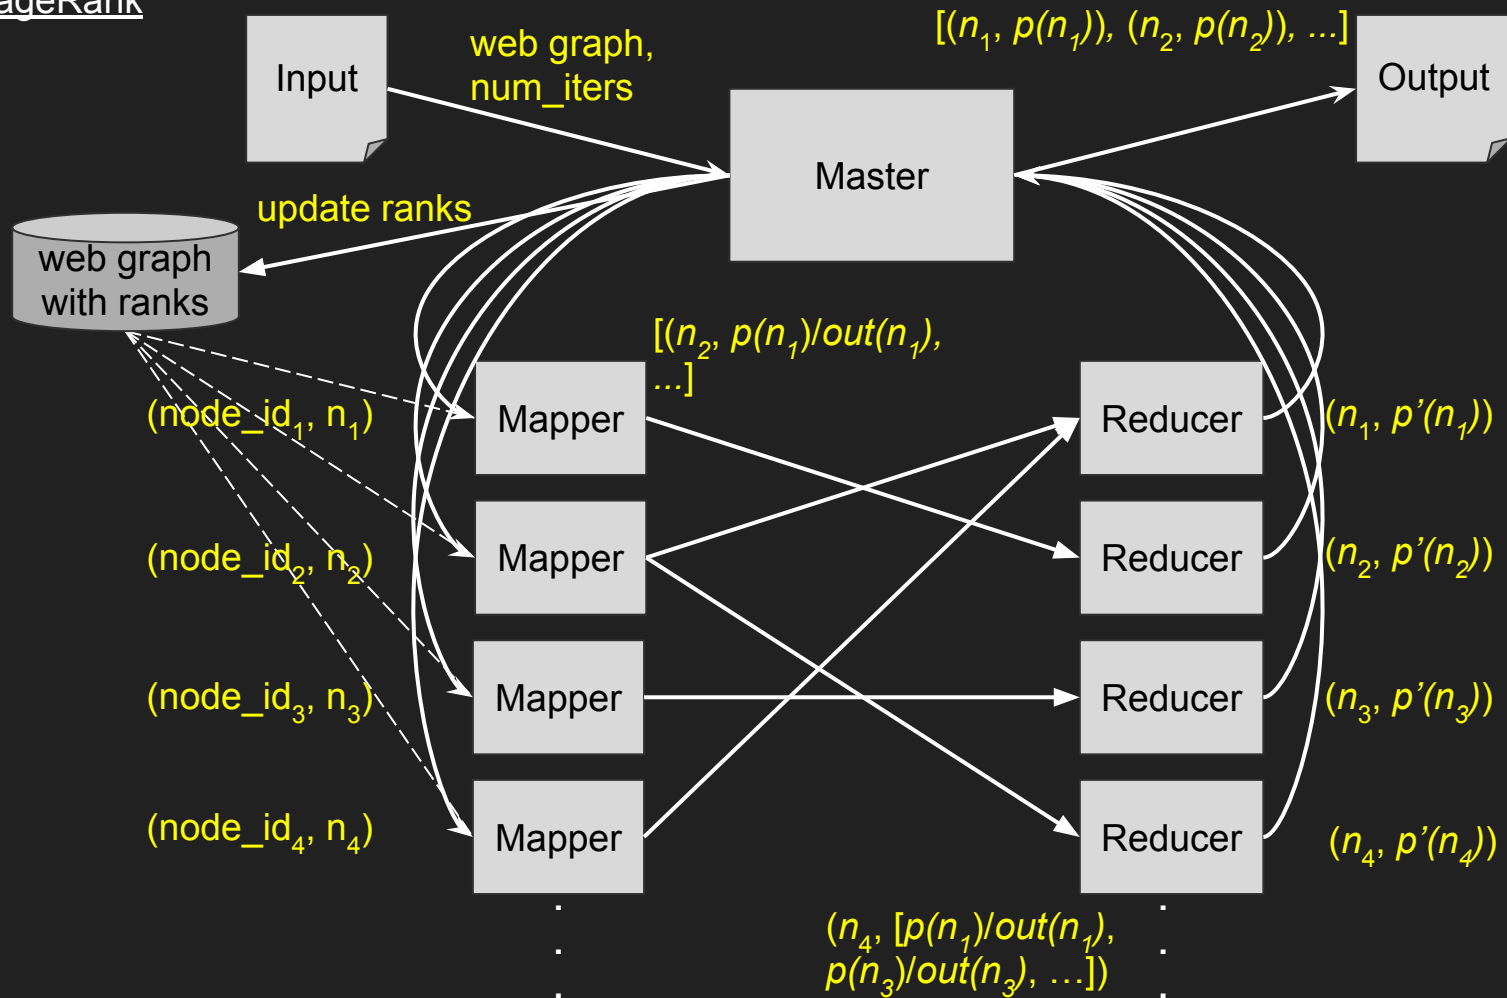
Inverted Index



k-means
clustering




PageRank




Consensus

Consensus Algorithms: Byzantine Generals




Consensus Algorithms: Byzantine Generals




Consensus Algorithms: Byzantine Generals

- Key Question: How do we coordinate with all the other generals at once?
- Assume we can't send signals the enemy can see (like torches)
- We're going to have to send messengers




Consensus Algorithms: Byzantine Generals

- What issues might we have?
 - How do we know the messengers made it?
 - How do we know that the message wasn't intercepted and replaced?
 - Same for the response
 - How do we know that the generals will even go along with the plan?




Consensus Algorithms: Byzantine Generals

- How does this relate to Computer Science?
 - Replace generals with computers and messengers with network packets



Consensus Algorithms: Byzantine Generals

- How does this relate to Computer Science?
 - Replace generals with computers and messengers with network packets



Consensus Algorithms: Byzantine Generals

- What issues might we have?
 - How do we know the packets made it?
 - Networking (TCP)
 - How do we know that the packet wasn't intercepted and replaced?
 - Encryption, Digital Signatures, etc.
 - How do we know that the other computers aren't working against us?

